

Yasin Saleem

yasinkoppara@gmail.com | [linkedin.com/in/yasinsaleem](https://www.linkedin.com/in/yasinsaleem) | github.com/YasinSaleem

Summary

AI/ML Engineer (B.Tech CSE, 2026) built and deployed production-grade RAG systems and multi-tenant AI infrastructure across healthcare and legal domains. Deployed citation-grounded LLM pipelines on AWS serving internal research workflows. **U.S. Citizen.**

Experience

AIML Engineer Intern, Healthcare-focused AI startup (Stealth Mode) – Remote December 2025 – Present

- Built Apache Airflow pipelines to ingest, preprocess (GROBID parsing), and chunk documents for automated knowledge base updates.
- Assisted in the design and development of a multi-tenant AI infrastructure isolating compute, storage, and network layers per entity, enabling secure deployment across healthcare clients.
- Led full production deployment (AWS EC2, Nginx, PM2, RDS).

AIML Intern, Turn2Law – Hybrid October 2025 – January 2026

- Developed FastAPI-based LLM agent to automate drafting of legal documents (Offer Letters, NDAs, MOUs) in <30 seconds.
- Integrated structured document templating using python-docx for formatting and compliance consistency.
- Contributed to RAG-based legal query system with citation-grounded responses.

AI Full Stack Developer Intern, Groweasy.ai – Remote July 2025 – October 2025

- Built TypeScript backend automation for TikTok campaign, ad group, and ad lifecycle management after analyzing 25+ API endpoints; reduced manual setup time by 30%.
- Improved campaign efficiency by 15% through structured API orchestration and workflow optimization.
- Prototyped AI-powered lead classification (WhatsApp chatbot) and dynamic website generation tools using Python to structure user data into SEO-optimized templates.

Research & Publications

Federated Learning for Thoracic Disease Detection, *Malque Publishing* [[link](#)] March 2025 – May 2025

- Benchmarked federated learning models on 10GB+ chest X-ray datasets.
- Proposed FedPA, a performance-based averaging strategy improving classification efficiency.
- Evaluated using F1, recall, precision, and accuracy metrics for privacy-preserving diagnostics.

Machine Learning in Healthcare, *Springer* [[link](#)] September 2024 – December 2024

- Improved treatment-seeking behavior prediction using XGBoost across six ML models.
- Benchmarked ensemble and kernel-based methods; evaluated via F1, recall, and accuracy.

Projects

AutoFT-Builder – Semantic Synthetic Dataset Pipeline

- Built CLI framework to generate synthetic fine-tuning datasets (1,000+ curated pairs) using LLM-based generation.
- Implemented embedding-based semantic deduplication (BGE-small + cosine similarity) to eliminate redundant samples.
- Designed SQLite-backed incremental storage, JSONL export, and configurable similarity thresholds.
- Achieved 98% test coverage (pytest) and optimized for rate-limited APIs.

Skills

Languages: Python, TypeScript, Go, SQL

LLM Systems: RAG pipelines, Hybrid retrieval (FTS + pgvector), Embeddings, Evaluation (RAGAS), Prompt design

Backend & Infra: FastAPI, Express, AWS (EC2, RDS, S3), Docker, Nginx, PM2, Apache Airflow

ML & Data: PyTorch, scikit-learn, TensorFlow, Pandas, NumPy, PySpark

Education

SRM Institute of Science and Technology – B.Tech in CSE (AI-ML) - **CGPA: 9.64/10.0**

Expected 2026